

Freight Rate Prediction

Machine Learning Engineer assessment. This report covers how I split the data, what exploration turned up, the data-quality problems I found and what I did about them, the model, and the December chart.

1. Summary

I predict `posted_rate` for all 12,000 validation loads with a two-stage model: a linear stage over smooth, extrapolating features, plus a gradient-boosted correction fitted to its residuals. The target is log rate-per-mile rather than the dollar rate.

On forward-in-time validation folds that reproduce the submission's forecast horizon, the model lands at **1.58% mean absolute percentage error** on the rows that are actually predictable, against a **1.35%** noise floor measured with a random split. Roughly 1.3% of rows carry a corrupted `posted_rate` that no model can recover; those are reported separately rather than being allowed to hide the rest.

The finding that mattered most is that the rate level is not flat over the calendar. It climbs through each quarter and resets at the quarter boundary, rising about 5% over a quarter's final month. December is a quarter-end month, so this drives the shape of the required chart.

2. How I split the data

The two files are separated in time, not at random:

File	Rows	Dates
<code>train_test.csv</code>	48,000	2025-01-01 to 2025-10-31
<code>validation.csv</code>	12,000	2025-11-01 to 2025-12-31

Validation is one to two months past the end of the labelled data. That rules out a random split for model selection. A random split lets the model see loads from the very days it is scoring, and the daily market level is the single largest source of variation in the target, so a random split would report an accuracy the submission cannot reproduce.

I used three forward-chaining folds instead. Each trains on everything before a cut date and scores a contiguous one-to-two-month block after it:

Fold	Trains on	Scores
1	before 2025-07-01	2025-07-01 to 2025-08-31
2	before 2025-08-01	2025-08-01 to 2025-09-30
3	before 2025-09-01	2025-09-01 to 2025-10-31

Fold 3 is the closest analogue to the real task: two months of unseen future, trained on eight months of history. I still ran a random split, but only to measure the noise floor. The gap between it and the temporal folds is what the forecast horizon costs, and reporting it as a validation score would be misleading.

One thing the split does use is the validation covariates. `market_index` is supplied for every validation row, and the daily average of it is the strongest single predictor of where rates sit on a given date. Using it is not leakage, since it is a feature and not the label, and it matches how the model would run in production, where today's market conditions are known. The folds are built the same way so that the validation numbers are honest about it.


3. What exploration turned up

Rate scales with distance, but less than proportionally. The elasticity of total rate to distance runs from about 0.85 at 300 miles to 0.89 at 2,500. Long hauls cost more in total and less per mile. This is why I model log rate-per-mile and carry log distance as a feature: it moves the dominant, nearly log-linear effect into a form the linear stage handles exactly, instead of asking a tree to approximate a smooth slope with steps.

Equipment carries a clean premium. Against Dry Van, Reefer runs about 12.9% higher per mile and Flatbed about 8.4%.


Geography is latitude, not city identity. Per-city effects correlate -0.84 with latitude: southern origins and destinations price higher per mile, northern ones lower, at roughly -0.34% per degree north on each end. Longitude does essentially nothing. This matters practically, because validation introduces eight cities that never appear in training.

The market level moves with the calendar. After dividing out the load mix, the daily rate level shows a clear sawtooth. It climbs through each quarter and drops at the quarter boundary. Underneath that sits a steady upward drift of about 0.60% per month.




Market rate level by day

Aligning the three complete quarters on days-to-quarter-end shows the same shape each time: flat until roughly 30 days out, then a climb of about 5%.



The quarter-end ramp

market_index is the main date-level driver. Its daily mean explains about 96% of the day-to-day movement in the rate level once the trend and the quarter ramp are accounted for. It also carries a weekly cycle, peaking Thursday and troughing Sunday, and rates follow it: about a 2.1% swing across the week. The weekday itself adds almost nothing once the index is in the model, so this is a market effect rather than a calendar effect.




Market index and the weekly cycle

quote_signal is noise. Its partial correlation with the target, after distance and the market level, is 0.03. Including it made held-out accuracy measurably worse (1.75% against 1.58% MAPE), because the booster fits the noise. I dropped it.

4. Data quality

Issue	Extent	What I did
Negative weight	292 rows	Took the absolute value. The negative entries' magnitudes match the positive distribution exactly, same 5,000 to 47,500 bounds and near-identical quartiles, so the sign is corruption and not a code for something.
Corrupted posted_rate	~1.3% of rows	Left in the file, excluded from fitting. They form two clean lobes at roughly 3.3x and 0.29x the true rate, well separated from the main body. The linear stage is fitted by iteratively reweighted least squares that trims beyond 4 robust standard deviations; 749 rows (1.56%) were dropped.

Issue	Extent	What I did
Missing weight	300 train, 165 validation	Imputed with the training median, plus an explicit missingness flag so the model can price the uncertainty.
Missing market_index	374 train, 249 validation	Filled from that date's mean. The daily level is what carries the signal, so a same-day fill loses very little.
Cities absent from training	8 cities, 1,447 validation rows (12.1%)	Used latitude and longitude rather than city identity. A city-name encoding would have no value for an eighth of the validation set; coordinates generalise, and latitude is what the city effect actually tracks.
distance floored at 70 miles	48 rows	Left alone. Short lanes are clipped to a 70-mile minimum, which looks like a billing floor rather than an error, and the model sees the same floor at prediction time.



Data quality

I also checked for duplicate `load_ids` and duplicate rows and found none. Each city maps to exactly one coordinate pair, and pickup and delivery coordinates agree for every city. The given distance correlates 0.9995 with great-circle distance between those coordinates, at a ratio of about 1.18, which is a normal road-circuitry factor. The coordinates are internally consistent even though they do not match real-world US geography, so I treated them as a valid synthetic geography and used them as-is.

5. Model

Target. Log rate-per-mile. Working in logs makes the multiplicative structure additive, which is what the data shows: equipment, market level and geography all act as percentage adjustments rather than dollar ones.

Stage one, linear. Least squares over log distance and its square, log weight, the daily market level and the within-day deviation from it, pickup and delivery coordinates, equipment indicators, a linear time index, and a hinge basis on days-to-quarter-end. Fitted robustly.

Stage two, gradient boosting. `HistGradientBoostingRegressor` fitted to the first stage's residuals, over the same features minus the time index.

Why split it this way. Validation runs two months past the end of training, and the rate level drifts upward the whole time. A tree cannot extrapolate that drift; past its last split it repeats its final value, which would leave every

November and December prediction short. A linear term in the time index does extrapolate it. So the linear stage carries everything smooth that has to run past the training window, and the booster only ever sees features whose values in November and December fall inside the range it was trained on. Days-to-quarter-end, day-of-week and day-of-month all satisfy that; raw month and day-of-year do not, and are deliberately absent.

Calibration. I predict the conditional median in log space and exponentiate, without a smearing correction. The corrupted rows inflate the mean of the actuals by about 1.3%, so scaling predictions up would reduce raw bias but would roughly double the median error. For any percentage or absolute error metric the uncorrected prediction is the right target.


6. Results

Forward-chaining folds. "Clean" columns exclude rows whose actual is more than 1.8x or less than 0.6x the prediction, which is the corrupted population.

Fold	Rows	MAE	RMSE	MAPE	MAE (clean)	RMSE (clean)	MAPE (clean)
1: Jul-Aug	9,671	92.77	618.41	4.05%	40.66	58.07	1.75%
2: Aug-Sep	9,429	84.84	611.36	3.74%	34.06	49.24	1.47%
3: Sep-Oct	9,523	91.85	629.80	4.05%	35.74	51.83	1.52%
Mean		89.82	619.86	3.95%	36.82	53.05	1.58%

Random-split reference, not a validation estimate: 1.35% MAPE (clean), 31.95 MAE (clean).


The gap between the columns is the corrupted 1.3%. They move RMSE from 53 to 620 while barely touching the clean error, which is why both are reported. The gap between 1.58% and the 1.35% floor is what two months of forecast horizon costs, and it is small, which says the calendar structure is being extrapolated correctly rather than guessed.



Held-out performance

Errors are close to symmetric in log space with a median ratio of 1.005, so the model is not systematically high or low. The final model was refitted on all 48,000 training rows.

7. December chart



December 2025 predicted rate

Predictions rise from **\$832.79 on 1 December to \$878.09 on 31 December**, about 5.4% across the month, or \$2.31 to \$2.44 per mile on a 360-mile Dry Van lane.

Two things drive the shape. The upward slope is the quarter-end ramp: December is the last month of Q4, and the same climb appears in the final month of Q1, Q2 and Q3. The weekly oscillation on top of it is the `market_index` cycle, peaking midweek and bottoming at the weekend. Both are patterns measured in the training data rather than artefacts. I checked the weekly wiggle specifically: the swing implied by the market index alone is about 2.1%, and the model's is close to that, so it is not the model amplifying noise.

Because the chart inputs carry no `market_index`, the daily level for each December date is taken from the validation file, which covers the same dates. The alternative was to hold the index at a constant and show only the trend and the ramp. I preferred using the observed daily level because it is the same information the model uses for the scored predictions, and a chart built on a different basis than the submission would not represent the model being submitted.

8. Limitations

- **The trend is extrapolated on ten months of history.** A linear fit held up best out of sample (degree-2 and degree-3 fits were worse on held-out months), but with under a year of data I cannot separate a genuine secular trend from the slow arm of an annual cycle. If the drift is actually seasonal and turns over in Q4, December predictions are biased high.
- **The quarter-end ramp is estimated from three quarters.** The shape is consistent across all three, but the Q4 ramp is an out-of-sample assumption, and Q4 could behave differently because of holiday freight.
- **Corrupted rows cap the achievable score.** At about 1.3% of rows with roughly 3.3x and 0.29x multipliers, they set a floor on RMSE no model can go below. If the scoring metric is RMSE on raw actuals, that term dominates and differences between reasonable models will be mostly invisible.
- **Eight validation cities are unseen.** Coordinates handle them, but Laredo sits at latitude 25.5, below the training range that starts at 28.4, so those rows extrapolate the latitude effect rather than interpolating it.

With more time I would fit the quarter-end ramp as a monotone shape-constrained spline rather than a hinge basis, and add quantile predictions so corrupted rows can be flagged rather than only trimmed.